

Charm production in muon deep-inelastic scattering with *FASERCal*

Yields, charm/anticharm tagging, and Bjorken- x reconstruction

Internal note — generated from the `muondis-fasercal` analysis

August 3, 2026

Abstract

We study charm production in high-energy muon deep-inelastic scattering (DIS) in *FASERCal*, using the POWHEG + Pythia8 muon-DIS samples developed for the *FASER ν* programme, and the *FASERCal* geometry and detector performance as given in the Conceptual Design Report and the July 2026 collaboration-meeting talk. Three results. First, the charm/anticharm asymmetry A_c — the observable that originally motivated the study — is *identically zero* in the available PDF sets, so it cannot be used as an intrinsic-charm observable without a PDF set carrying an explicit charm asymmetry. Second, the accessible intrinsic-charm signal is instead the large- x_B charm excess, and we confirm that the lepton-only x_B reconstruction used in the literature is destroyed by *FASERCal*’s muon-momentum resolution. Third, and centrally, we show that *calorimetric* reconstruction methods — which *FASER ν* cannot use — recover a factor ~ 3 – 3.8 of the large- x_B signal, and raise the intrinsic-charm purity of the selected sample from 64% to 95% when two statistically independent methods are combined. All uncertainties quoted are statistical and detector-resolution effects only; systematic uncertainties are discussed in Sec. 6 but are *not* included in any number.

Contents

1	Motivation and scope	1
2	Inputs and their provenance	2
2.1	Monte Carlo samples	2
3	The muon-DIS charm chain	2
3.1	Charm tagging without a vertex	2
3.2	Yields	3
3.3	Absorber and AHCAL options	3
4	The charm/anticharm asymmetry is not a usable observable	3
4.1	It is zero by construction	3
4.2	But fragmentation generates one	4
4.3	The signal that does exist	4
5	Bjorken-x reconstruction	4
5.1	Kinematics and notation	4
5.2	Method 1: lepton-only	5
5.3	Method 2: Jacquet–Blondel	5
5.4	Method 3: the Σ method	6
5.5	Method 4: double-angle	6
5.6	Validation	7

5.7	Results	7
5.8	Why combining works	7
5.9	Intrinsic-charm content of the selection	7
6	Systematic uncertainties — NOT included	8
6.1	What was varied, and what it does	8
6.2	An actionable consequence: tracker granularity	9
6.3	Systematics not evaluated at all	9
6.4	Statistical limitations	10
7	A methodological problem worth recording	10
8	Summary	10

1 Motivation and scope

The proton may carry a non-perturbative, “intrinsic” charm (IC) component: a $|uudc\bar{c}\rangle$ Fock state in which the charm quarks carry a large fraction of the proton momentum. Its signature is an excess of charm at large Bjorken x_B over what perturbative gluon splitting produces. High-energy muon DIS at the LHC is a clean probe, and *FASERCal* — an electronic, highly segmented three-dimensional scintillator calorimeter with a downstream magnetised muon spectrometer — is a candidate detector.

This note addresses three questions:

1. How many muon-DIS charm events does *FASERCal* collect, and can charm be distinguished from anticharm? (Sec. 3)
2. Is the charm/anticharm asymmetry A_c a usable intrinsic-charm observable? (Sec. 4)
3. Can x_B be reconstructed well enough at large x_B for the intrinsic-charm excess to survive? (Sec. 5)

The third is the critical one. Ref. [1] reconstructs x_B from the scattered muon alone and finds that a muon-momentum smearing of 10% already reduces the $x_B \geq 0.2$ charm excess from $\times 8$ to $\times 2$, and that 30% removes it entirely. *FASERCal*’s spectrometer is iron-core and multiple-scattering dominated, so this is a genuine threat — but *FASERCal* is also a calorimeter, which *FASER ν* is not.

2 Inputs and their provenance

Every number used is listed in Table 1 with its source. Where a quantity is an assumption rather than a documented value, this is stated explicitly; those are the entries that most need review.

2.1 Monte Carlo samples

Two POWHEG + Pythia8 productions differing *only* in the charm PDF:

Production	LHAPDF	Role
<code>production_v1</code>	331100, NNPDF40_nnlo_as_01180	fitted charm (IC allowed)
<code>production_pc</code>	332100, NNPDF40_nnlo_pch_as_01180	perturbative charm (null)

Muon beam on a fixed nucleon target, $Q_{\min} = 1.65$ GeV, four per-nucleon samples ($\mu^\pm$ on p/n). 25 independent seeds, 1.90×10^6 showered events per hypothesis. The muon flux is implemented as a beam PDF of a fictitious 7 TeV beam [1]; this has consequences discussed in Sec. 7.

Target composition is applied at analysis time: *FASERCal* is polystyrene (CH, $w_p = 0.538$) plus tungsten ($w_p = 0.402$), not the tungsten-only target of *FASER ν* .

Quantity	Value	Source
Run 4 luminosity	780 fb^{-1}	FASER <i>Cal</i> talk; CDR Sec. 3.1
3DCal geometry	10 modules $\times$ 20 layers of 1 cm^3 cubes	talk slide 4
3DCal transverse face	$48 \times 48 \text{ cm}^2$	talk slide 11
Tungsten absorber	1, 5 or 10 mm <i>per module</i>	talk slide 39; CDR Table 5
3DCal total mass	581 / 896 / 1118 kg	CDR Table 5
Fiducial region	$z < 1150 \text{ mm}$ (48% of length)	CDR Table 8
Position	LoS shift (452, 236) mm, tilt 4.5°	talk slides 11–12
Muon momentum resolution	20/22.8/29/63.1% at 20/100/200/1000 GeV	CDR Sec. 2.6.2
Charge misidentification	$< 0.7\%$ below 100 GeV; 2–5% at TeV	CDR Sec. 2.6.2
Spectrometer acceptance	43% (≥ 2 MDT stations)	talk slide 33
Hadronic energy resolution	$\sim 9\%$	talk (p_{jet} , NC)
Off-axis flux factor	$F_{\text{flux}} = 0.96$	<i>this work</i> , from FLUKA
Run 4 optics enhancement	$F_{\text{optics}} = 2.0$	CDR Sec. 1.3.2 (rate $\times 2$)
Muon angular resolution	0.3 mrad	<i>derived</i> from 3DCal geometry
Hadronic angular resolution	20 mrad	ASSUMPTION – no source
Identification + linking eff.	scanned 0.05–1	not assumed; scanned

Table 1: Inputs. Entries above the rule are taken from FASER*Cal* documents; below the rule are derived or assumed in this work. Where the talk and the CDR disagree, the talk (more recent) is used.

3 The muon-DIS charm chain

3.1 Charm tagging without a vertex

In emulsion, charm is identified by imaging the sub-mm decay kink. In a scintillator this is impossible: the charm flight length is a few mm, at or below cell size, and in the interleaved-tungsten configuration the decay happens inside an absorber plate. **There is no displaced-vertex charm tag in FASER*Cal*.**

The surviving handle is the muon from the semileptonic charm decay, whose charge *is* the charm sign:

$$c \rightarrow s W^+, W^+ \rightarrow \mu^+ \nu \quad \bar{c} \rightarrow \bar{s} W^-, W^- \rightarrow \mu^- \bar{\nu}. \quad (1)$$

This gives an opposite-sign dimuon: the scattered DIS muon plus a softer muon of opposite charge. FASER*Cal* itself has no field; the sign comes from the downstream spectrometer. The identity $\text{sign}(c) = q_\mu$ is verified event-by-event in the code and holds for 100% of events.

3.2 Yields

Table 2 gives the chain at Run 4 nominal conditions.

Stage	3DCal only	+AHCAL	Fraction
Muon DIS in fiducial volume	373 601	2 503 837	100%
containing charm	10 578	71 489	2.86%
charm $\rightarrow$ semileptonic μ	1 866	12 613	17.6% of charm
μ punches out of calorimeter	1 523	10 301	81.7%
reaches spectrometer (43%)	666	4 509	43%
identified + linked ($\varepsilon = 0.9$)	599	4 058	0.16% of DIS
Charge-tag purity	99.1%		CDR Sec. 2.6.2
Statistical reach $\sigma(A_c)$	0.041	0.016	

Table 2: Muon-DIS charm chain, Run 4 (780 fb^{-1}), 3DCal with 1 mm W per module, fiducial $z < 1150 \text{ mm}$ (277 kg), $F_{\text{flux}} = 0.96$, $F_{\text{optics}} = 2.0$.

Two features dominate. The chain is **branching-limited, not detector-limited**: $2.86\% \times 17.6\% \approx 1/200$ before any detector effect, and no instrumentation improvement recovers it. And 94% of charm events contain *two* charm hadrons — the signature of photon-gluon fusion $\gamma g \rightarrow c\bar{c}$ — which is why the per-event semileptonic rate (17.6%) is about twice the per-hadron branching ratio.

Acceptance is an *angular* problem, not a momentum one: the decay muon has a median angle of 33 mrad against the scattered muon’s 4.3 mrad, and 83% already pass the punch-through momentum threshold. The assumed 43% acceptance corresponds to a $\theta < 8.7$ mrad cone. Charge identification is *not* a bottleneck: with a median decay-muon momentum of 11.5 GeV, the CDR misidentification is $< 0.7\%$.

3.3 Absorber and AHCAL options

Configuration	Fiducial mass	Tagged	$\sigma(A_c)$
1 mm W / module (baseline)	277 kg	599	0.041
5 mm W / module	421 kg	901	0.033
10 mm W / module	514 kg	1102	0.030
+ AHCAL (1.60 t fiducial, $\sim 98\%$ Fe)	1 877 kg	4058	0.016

The AHCAL adds $5.8\times$ the 3DCal target mass. Its $4 \times 4 \text{ cm}^2$ granularity makes vertex finding and muon linking harder, so its curve should be read at a *lower* identification efficiency — which is why the main result is presented against ε rather than at a fixed value.

4 The charm/anticharm asymmetry is not a usable observable

4.1 It is zero by construction

Direct inspection of the LHAPDF grids gives, at $Q = 10 \text{ GeV}$:

x	$xc(x)$	$x\bar{c}(x)$	$(c - \bar{c})/(c + \bar{c})$
0.01	2.043×10^{-1}	2.047×10^{-1}	-1.0×10^{-3}
0.30	8.635×10^{-3}	8.669×10^{-3}	-2.0×10^{-3}
0.70	5.190×10^{-4}	5.191×10^{-4}	-7.3×10^{-5}

Over all 2352 grid points the maximum of $|c - \bar{c}|/(c + \bar{c})$ is 6.4×10^{-3} , and the number sum rule gives $\int (c - \bar{c}) dx / \int (c + \bar{c}) dx = -1.5 \times 10^{-4}$. NNPDF4.0 fits the total $c + \bar{c}$ and sets $c - \bar{c} = 0$. The residual is *larger* in the perturbative set than the fitted one — the signature of numerical noise rather than physics.

This is consistent with the light-cone argument: the minimal $|uudc\bar{c}\rangle$ Fock state contains one c and one $\bar{c}$ with identical distributions, so BHPS intrinsic charm predicts $A_c(x) = 0$ at every x . A non-zero A_c requires physics beyond the minimal Fock state, is model-dependent, and is constrained by $\int (c - \bar{c}) dx = 0$.

4.2 But fragmentation generates one

Even with an exactly symmetric charm PDF, the *observable* sample is asymmetric, through the leading-baryon effect:

Species	$N(c)$	$N(\bar{c})$	Ratio	BR($\rightarrow \mu$)
Λ_c	13 430	1 963	6.84	4.8%
D^+	16 683	20 355	0.82	16.7%
D^0	31 730	38 820	0.82	6.6%

The c quark can pick up two valence quarks from the target remnant to form a Λ_c ; the $\bar{c}$ cannot easily form an antibaryon. Λ_c also has the lowest semileptonic branching ratio of any charm hadron, so the c is preferentially *hidden from the muon tag* while the $\bar{c}$ is not, leaving a $\bar{c}$ -enriched tagged sample: $A = -0.061 \pm 0.036$ ($N_{\text{eff}} = 778$, i.e. 1.7σ).

Consequences. A_c cannot be used as an IC observable with existing PDF inputs, and the fragmentation-induced asymmetry is of the *same size* as our statistical reach (0.041), so it is a competing background rather than a correction. Attention should move to the large- x_B excess.

4.3 The signal that does exist

The fitted/perturbative charm ratio at $Q = 10$ GeV:

x	0.2	0.3	0.4	0.5	0.7
fitted / perturbative	1.8	3.2	5.7	10.9	44.3

At *event* level, integrated above $x_B = 0.2$, the charm sample is $\approx 99\%$ IC-attributable (Sec. 5.9).

5 Bjorken- x reconstruction

5.1 Kinematics and notation

For fixed-target DIS with nucleon mass M , incoming muon energy E_μ , scattered muon energy $E_{\mu'}$ and scattering angle θ_μ :

$$\nu = E_\mu - E_{\mu'}, \quad Q^2 = 2E_\mu E_{\mu'}(1 - \cos \theta_\mu), \quad x_B = \frac{Q^2}{2M\nu}, \quad y = \frac{\nu}{E_\mu}. \quad (2)$$

The hadronic system recoiling against the muon is exactly $X = q + p_{\text{target}}$, so it is fully determined by the muon kinematics:

$$E_{\text{had}} = \nu + M, \quad p_T^{\text{had}} = E_{\mu'} \sin \theta_\mu, \quad p_z^{\text{had}} = |p_\mu| - E_{\mu'} \cos \theta_\mu. \quad (3)$$

The key numbers of the problem, from our sample:

Median E_μ	665 GeV
Median $E_{\mu'}$	469 GeV
Median ν	52 GeV
Median y	0.161
Median θ_μ	4.3 mrad
Median θ_h	81 mrad

ν is only 7.7% of E_μ . It is a small difference of two large numbers, and this single fact drives everything that follows.

5.2 Method 1: lepton-only

What it does. Uses only the muon. Both muon energies and the scattering angle are measured; ν and Q^2 follow directly:

$$\nu^{\text{lep}} = E_\mu - E_{\mu'}, \quad Q_{\text{lep}}^2 = 2E_\mu E_{\mu'}(1 - \cos \theta_\mu), \quad x_{\text{lep}} = \frac{Q_{\text{lep}}^2}{2M\nu^{\text{lep}}}. \quad (4)$$

Why it is used. It is the only method available to an emulsion detector. FASER ν measures angles superbly and energy not at all, so it is *forced* into this choice. It is what Ref. [1] uses.

Why it fails here. Because ν is a difference of two large numbers, a fractional error ϵ on each muon leg propagates as

$$\frac{\sigma_\nu}{\nu} \approx \frac{\epsilon}{y} \sqrt{1 + (1 - y)^2} \sim \frac{\epsilon}{y}. \quad (5)$$

With $y = 0.16$ and the CDR resolution $\sigma_p/p \approx 38\%$ at 469 GeV, this gives a *several-hundred-percent* error on ν . Measured directly, the ν resolution is an interquartile range of 410%. The migration matrix (Fig. 1) shows essentially *no* correlation between x_{true} and x_{reco} .

Additional problem. It requires E_μ . The downstream spectrometer sees only the *outgoing* muon, so if the incoming muon is not measured elsewhere this method **does not exist at all**.

5.3 Method 2: Jacquet–Blondel

What it does. Uses only the *hadronic* system. The energy transfer is measured directly by the calorimeter rather than inferred from a difference:

$$\nu^{\text{JB}} = E_{\text{had}} - M, \quad Q_{\text{JB}}^2 = \frac{(p_T^{\text{had}})^2}{1 - y}, \quad x_{\text{JB}} = \frac{Q_{\text{JB}}^2}{2M\nu^{\text{JB}}}. \quad (6)$$

Why the target-mass subtraction matters. The measured hadronic energy is $\nu + M$ (Eq. 3); it includes the struck nucleon’s rest energy. Omitting M biases x_{JB} low by M/ν , which is 9% on the $x_{\text{B}} \geq 0.2$ sample where $\nu \approx 9$ GeV. This is a small but real effect that is easy to miss.

Strengths. The muon momentum is not used *at all*: a 9% hadronic resolution gives a 9% error on ν , not an amplified one. E_μ enters only through the mild $(1 - y)$ factor, so the method survives when E_μ is unmeasured.

Weaknesses. Q^2 comes from the hadronic transverse momentum, which is the *difference* of a nearly-forward system — so it is sensitive to the hadronic *angular* resolution. In practice JB has excellent efficiency (78%) but poor purity (34%): the vastly more numerous low- x_{B} events migrate upward and flood the selection.

5.4 Method 3: the Σ method

The idea. Reconstruct the *incoming* muon energy from the final state, so that it need not be measured. With the beam along $+z$, the initial state has $E + p_z = 2E_\mu + M$, and since $E + p_z$ is conserved,

$$2E_\mu = \sum_{\text{had}} (E + p_z) + E_{\mu'}(1 + \cos \theta_\mu) - M, \quad (7)$$

after which one proceeds exactly as in the lepton-only method but with this E_μ .

A necessary derivation. The collider Σ method uses $E - p_z$ [4]. For a *fixed target* with the beam along $+z$ that combination is identically M for the initial state and therefore carries no information. The correct fixed-target analogue uses $E + p_z$, as in Eq. 7. This was derived for this work.

Why it works — an exact cancellation. At the small angles relevant here, $E + p_z \approx 2E$, and Eq. 7 reduces to simple energy conservation, $E_\mu \approx E_{\mu'} + E_{\text{had}} - M$. Then

$$\boxed{\nu^\Sigma = E_\mu^{\text{reco}} - E_{\mu'}^{\text{reco}} = (E_{\mu'} + E_{\text{had}} - M) - E_{\mu'} = E_{\text{had}} - M} \quad (8)$$

The muon-momentum error appears in both terms and cancels exactly. The resulting ν resolution is 13% (interquartile), against 410% for lepton-only where E_μ and $E_{\mu'}$ are measured

independently and their errors *add*. This is an algebraic cancellation, not a tuned improvement, which is why the method is robust.

Insensitivity to the hadronic angle. $p_z = |p| \cos \theta_h$, and with $\theta_h \approx 81 \text{ mrad}$ a 20 mrad smearing changes $\cos \theta_h$ by 0.2%. The method is therefore *first-order insensitive* to the hadronic angular resolution — the single largest unknown in the calorimetric methods.

Practical note. A user asking “can we measure E_μ by taking the spectrometer muon and adding the energy that went into the interaction?” is describing exactly this method. The two constructions are numerically identical here.

5.5 Method 4: double-angle

What it does. Uses only the two *angles*, plus the beam energy. Solving energy–momentum conservation for the scattered muon energy:

$$E_{\mu'}^{\text{DA}} = E_\mu \frac{\sin \theta_h}{\sin(\theta_\mu + \theta_h)}, \quad (9)$$

then Q^2 and ν as usual. Being built from angles, it is immune to *energy-scale* errors in both the calorimeter and the spectrometer — a genuine advantage, since energy scales are the hardest thing to calibrate.

Weaknesses. Two.

1. It requires E_μ , and $x_{\text{DA}} \propto E_\mu$ linearly. If the incoming muon is not measured, the method **does not exist**.
2. It depends entirely on the hadronic angle θ_h , which in a calorimeter is a *shower axis*, not a track. Its performance degrades steeply with σ_{θ_h} (Fig. 2, left): the figure of merit falls from 0.65 to 0.36 between 0 and 50 mrad.

Since σ_{θ_h} is currently an assumption with no documented value, the double-angle numbers carry the largest modelling uncertainty of the four.

5.6 Validation

With all resolutions set to zero, all four methods return $x_{\text{reco}}/x_{\text{true}} = 1.0000$. The framework is therefore verified *before* any smearing is applied. This test was essential: it exposed both the missing target-mass term in JB and a serious problem with the hadronic energy definition (Sec. 7).

5.7 Results

The figure of merit is efficiency $\times$ purity for the $x_B \geq 0.2$ selection on charm events, where *efficiency* is the fraction of true $x_B \geq 0.2$ events reconstructed above the cut, and *purity* is the fraction of selected events that truly have $x_B \geq 0.2$.

5.8 Why combining works

The methods fail for *different* reasons, so their errors are not all correlated. Correlation of $\log(x_{\text{reco}}/x_{\text{true}})$:

Method	Efficiency	Purity	Eff.×Pur.	vs lepton	IC fraction
lepton-only	23.7%	63.8%	0.151	—	82%
Jacquet–Blondel	78.2%	33.6%	0.262	×1.7	38%
Σ	77.1%	57.6%	0.444	×2.9	64%
double-angle	82.5%	61.3%	0.506	×3.4	72%
Σ AND double-angle	66.3%	85.7%	0.568	× 3.8	95%
all three	54.9%	94.7%	0.520	×3.4	97%

Table 3: Large- x_B signal retention, $x_B \geq 0.2$, 25 seeds, all modelled effects included. “IC fraction” is the intrinsic-charm-attributable fraction of the selected sample, $1 - N_{\text{pert}}/N_{\text{fitted}}$.

	lepton	JB	Σ	DA
lepton-only	1.00	−0.03	0.65	−0.05
Jacquet–Blondel	−0.03	1.00	−0.01	0.94
Σ	0.65	−0.01	1.00	0.08
double-angle	−0.05	0.94	0.08	1.00

Σ and double-angle are **nearly independent** (0.08): one is driven by the calorimetric *energy*, the other by the *angles*. JB and double-angle are redundant (0.94), both being hadronic-system-driven. Requiring both members of the independent pair raises purity from 58% to 86% and the IC fraction of the selected sample from 64% to 95%.

Adding JB buys further purity but costs more efficiency than it gains. Note that this is a crude AND of two cuts; a proper constrained kinematic fit using all measurements with their covariances should do better still.

If E_μ is not measured. Double-angle and lepton-only cease to exist. The fallback is Σ AND Jacquet–Blondel (correlation −0.01, also independent): efficiency 60.8%, purity 84.5%, product 0.514 — still ×3.4 over lepton-only. Σ alone is unchanged at 0.444.

5.9 Intrinsic-charm content of the selection

At truth level the $x_B \geq 0.2$ cut is well placed:

x bin	0.1–0.2	0.2–0.3	0.3–0.4	> 0.4
IC fraction	85%	98%	99%	~ 100%

Integrated above 0.2 it is $\approx 99\%$. What is actually *selected* is diluted by migration from low x_B , where there is no IC — this is the purity restated in physics terms, and it is why the method combination matters (Table 3, last column).

Caveat: the perturbative sample above $x_B = 0.2$ has $N_{\text{eff}} = 7$ (966 raw events, one carrying 9% of the weight). The 99% survives only because the perturbative contribution is small to begin with; the perturbative *yield* itself must not be quoted, and any ratio-based statement there needs a dedicated high-statistics `production.pc`.

6 Systematic uncertainties — NOT included

Every number in this note is a statistical or detector-resolution effect. No systematic uncertainty is included anywhere. This section lists what is missing, with an estimate where one can be made cheaply.




Figure 1: Truth- $x_B \rightarrow$ reco- x_B migration for charm events. Dashed line is perfect reconstruction. Lepton-only shows almost no correlation; Σ retains a clear diagonal.




Figure 2: Left: signal retention versus the hadronic angular resolution. Σ is flat; double-angle degrades steeply. Right: versus whether the incoming muon energy is measured — lepton-only and double-angle do not exist without it.

Variation	Σ	double-angle	lepton-only
nominal	0.444	0.506	0.151
hadronic resolution 9% $\rightarrow$ 15%	0.442	0.506	0.151
hadronic resolution 9% $\rightarrow$ 5%	0.444	0.506	0.151
muon resolution $\times 1.3$	0.378	0.456	0.131
muon resolution $\times 0.8$	0.521	0.561	0.175
hadronic angle 20 $\rightarrow$ 50 mrad	0.444	0.356	0.151
no neutrino subtraction	0.433	0.506	0.151

Table 4: Sensitivity of the figure of merit to the modelled resolutions.

6.1 What was varied, and what it does

The dominant sensitivity is to the **muon momentum resolution**, even for the calorimetric methods — because although ν is calorimetric, Q^2 still uses the muon. The hadronic *energy* resolution is essentially irrelevant over 5–15%. The hadronic *angular* resolution matters only for double-angle.

6.2 An actionable consequence: tracker granularity

The CDR quotes the effect of finer SciFi granularity. Scaling to our momentum:

SciFi pitch	σ_p/p at 469 GeV	Σ	double-angle	Σ AND DA
100 μm (baseline)	38%	0.444	0.506	0.568
70 μm	29%	0.538	0.572	0.635
50 μm	26%	0.570	0.605	0.658

Going from 100 to 50 μm improves the combined selection by 16%. Since the muon resolution is the dominant lever, this is the single most effective detector change for this measurement.

6.3 Systematics not evaluated at all

1. **PDF and QCD-scale uncertainties.** The LHE files carry `rep001--rep103` NNPDF4.0 replica and scale weights, but only the nominal weight is read; the replicas are discarded at both the parsing and showering stages. *No PDF error band exists anywhere in this note.* For the fitted-vs-perturbative comparison this matters most, since the charm PDF at large x is precisely the poorly-known quantity.
2. **Absolute flux normalisation.** A residual factor ~ 1.7 discrepancy with Ref. [1] remains unexplained after correcting the reference target mass. $F_{\text{optics}} = 2.0$ is the CDR’s own estimate for a 2024 configuration, not Run 4, and the CDR states it is unclear whether the increase can be mitigated. $F_{\text{flux}} = 0.96$ assumes the FLUKA x convention matches the *FASERCal* one.
3. **Energy scale.** Only resolutions were varied, never a *bias*. A hadronic energy scale offset propagates directly into ν and hence x_B , and on a steeply falling spectrum a few-percent scale error can move the $x_B \geq 0.2$ yield substantially. This should be evaluated.
4. **Hadronisation model.** The $\Lambda_c/\bar{\Lambda}_c = 6.8$ leading-baryon effect, which drives the fragmentation asymmetry of Sec. 4, is Pythia-specific and is not varied.
5. **Nuclear effects.** Free-nucleon PDFs are used; no nuclear modification is applied for carbon or tungsten.
6. **Detector modelling.** No reconstruction simulation: no voxel clustering, no shower overlap, no pile-up from the $\sim 1 \text{ Hz/cm}^2$ muon background, and no π/K decay-in-flight background to the soft second muon.
7. **The hadronic angular resolution** $\sigma_{\theta_h} = 20 \text{ mrad}$ has no documented source.

6.4 Statistical limitations

The perturbative sample above $x_B = 0.2$ has $N_{\text{eff}} = 7$. Any statement involving the perturbative *yield* or a fitted/perturbative *ratio* above $x_B = 0.2$ is statistics-limited. The efficiency and purity numbers, which use only the fitted sample (7566 raw events above the cut, $N_{\text{eff}} = 5319$), are not affected.

7 A methodological problem worth recording

The hadronic four-vector is computed from momentum conservation (Eq. 3), *not* by summing Pythia final-state particles. The particle sum is contaminated: the muon flux is implemented as a beam PDF of a fictitious 7 TeV beam, and that beam’s remnant deposits a few GeV of hadronic activity that is not part of the DIS vertex. Energy conservation $E_\mu + M - E_{\mu'} - E_{\text{had}} \geq 0$ fails in 80% of events, rising to 98% on the $x_B \geq 0.2$ charm sample where E_{had}/ν reaches 2.5. It is negligible at large ν ($E_{\text{had}}/\nu = 1.02$ above 500 GeV) but dominates exactly where the physics is.

Three particle-level fixes were tried and all failed: dropping beam-remnant particles (their hadronic *descendants* remain); a forward-cone cut (moves the mean, not the median); and requiring mother-chain ancestry to reach the target beam (*worse*, $E_{\text{had}}/\nu = 8.5$, because Pythia’s colour reconnection leaves the ancestry graph too connected to separate the two beam sides). Using conservation is exact and makes the closure exact by construction.

The cost is that neutrino losses must be added back explicitly, which is done: neutrinos from charm semileptonic decays are nonzero in 35% of charm events and carry a median 13.2 GeV (8.4% of ν) when present. Their effect is -10% on JB and negligible on Σ .

8 Summary

Next steps

1. Evaluate PDF and scale uncertainties using the `rep001--rep103` weights, which requires carrying them through the showering step.
2. Evaluate energy-*scale* systematics, not just resolutions.
3. Replace $\sigma_{\theta_h} = 20$ mrad with a value from the FASER*Cal* reconstruction.
4. Implement a constrained kinematic fit (cf. Ref. [1] Sec. 4.2) in place of the crude AND of two cuts.
5. Generate a high-statistics perturbative-charm sample so that fitted/perturbative ratios above $x_B = 0.2$ become meaningful.
6. Simulate the π/K decay-in-flight background to the soft second muon.

References

- [1] R. Francener *et al.*, “Deep-Inelastic Scattering at TeV Energies with LHC Muons”, [arXiv:2506.13889](#).
- [2] FASER*Cal* Collaboration, “The FASER*Cal* Conceptual Design Report”, CERN-FASER-NOTE-2026-004, v0, 16 February 2026.
- [3] “FASER*Cal*”, FASER Collaboration Meeting, Bern, 15 July 2026.
- [4] U. Bassler and G. Bernardi, “On the kinematic reconstruction of deep inelastic scattering at HERA: the Σ method”, Nucl. Instrum. Meth. A **361** (1995) 197.

Question	Answer
How many tagged charm events?	599 charge-signed, vertex-linked charm muons in the 3DCal (4058 including the AHCAL) at 780 fb^{-1} .
Is charm/anticharm separable?	Yes in principle — the decay-muon charge <i>is</i> the charm sign, and the CDR spectrometer misidentifies $< 0.7\%$ — but the reach is $\sigma(A_c) = 0.041$.
Is A_c an intrinsic-charm observable?	No. It is identically zero in both available PDF sets. A non-zero A_c needs a PDF set with explicit charm asymmetry, or a model. A fragmentation-induced asymmetry of the same size as our reach competes with it.
What is the accessible IC signal?	The large- x_B charm excess: $\times 1.8$ at $x_B = 0.2$ rising to $\times 44$ at 0.7 . Above $x_B = 0.2$ the charm is $\approx 99\%$ IC-attributable.
Does muon smearing destroy it?	For lepton-only, yes — ν is 7.7% of E_μ , so 38% muon resolution gives a several-hundred-percent ν error and no correlation between x_{true} and x_{reco} .
Do calorimetric methods rescue it?	Yes. Σ retains $\times 2.9$ more signal; Σ AND double-angle $\times 3.8$, lifting the IC purity of the selection from 64% to 95% .
Which method should be used?	Σ , combined with double-angle where possible. Σ is flat in the hadronic angular resolution, unaffected by the neutrino loss, and the only method that works when E_μ is unmeasured.
Do neutral particles undermine it?	No. Neutrinos cost JB 10% and Σ nothing; neutral hadrons are 8% of the visible energy.
What most improves the measurement?	The muon momentum resolution . Going from 100 to $50 \mu\text{m}$ SciFi pitch improves the combined selection by 16% .
Are systematics included?	No. Nothing in this note includes a systematic uncertainty (Sec. 6). PDF uncertainties in particular are entirely absent, although the replica weights exist in the LHE files.

Table 5: Summary of findings.